

# Project - High Level Design

on

## Healthcare Planning Assistant Agent

Course Name: Agentic AI

**Institution Name:** Medicaps University – Datagami Skill Based Course

*Student Name(s) & Enrolment Number(s):*

Sr no	Student Name	Enrolment Number
01	MEET SOLANKI	EN22CS301595
02	MAYURI PATIDAR	EN22CS301593
03	MOKSHIKA YADAV	EN22CS301614
04	PALAK SHAHDADPURI	EN22CS301677
05	NIRDESH SHARMA	EN22CS301655
06	PANKAJ	EN22CS301679

*Group Name: 11D6*

*Project Number: AAI-22*

*Industry Mentor Name:*

*University Mentor Name: Prof. Nishant Srivastava*

*Academic Year: 2025-26*

# Table of Contents

1. Introduction.
  - 1.1. Scope of the document.
  - 1.2. Intended Audience
  - 1.3. System overview.
2. System Design.
  - 2.1. Application Design
  - 2.2. Process Flow.
  - 2.3. Information Flow.
  - 2.4. Components Design
  - 2.5. Key Design Considerations
  - 2.6. API Catalogue.
3. Data Design.
  - 3.1. Data Model
  - 3.2. Data Access Mechanism
  - 3.3. Data Retention Policies
  - 3.4. Data Migration
4. Interfaces
5. State and Session Management
6. Caching
7. Non-Functional Requirements
  - 7.1. Security Aspects
  - 7.2. Performance Aspects
8. References

# 1. Introduction

## 1.1 Scope of the Document

This document provides the High-Level Design (HLD) for the Healthcare Planning Assistant, an AI-powered healthcare advisory application.

The scope of this document includes:

- Overall system architecture and design approach
- Major modules and components of the system
- Workflow and data flow across system layers
- Integration with external AI services (Google Gemini API and Groq LLM API)

## 1.2 Intended Audience

This document is intended for:

- **Project evaluators and faculty members** reviewing the system design
- **Developers** who will implement, test, or enhance the application
- **System designers** responsible for architecture and modular design

## 1.3 System Overview

The Healthcare Planning Assistant is a web-based AI application designed to provide structured health guidance based on user-entered symptoms. The system uses Large Language Models (LLMs) through Google Gemini API and Groq LLM API to generate responses in a step-by-step manner.

The system provides:

- Possible health conditions (advisory, not diagnostic)
- Precautionary measures
- General medication guidance (non-prescriptive)
- Nutrition and diet recommendations

The application is developed using Python, Streamlit for the user interface, and LangChain for prompt orchestration and sequential chaining. The system is currently designed as a prototype for local execution and does not store user medical data permanently.

## 2. System Design

### 2.1 Application Design


The Healthcare Planning Assistant is designed as a modular AI-driven healthcare advisory application. The application follows a layered design to separate user interaction, business logic, LLM orchestration, and external service integration.

#### Layered Structure

- **Presentation Layer (Streamlit UI):** Handles user interaction, symptom input, model selection, and output display.
- **Application Logic Layer (Python Core):** Controls workflow execution, manages session state, and structures responses.
- **LLM Orchestration Layer (LangChain):** Manages prompt templates, sequential chaining, and model abstraction.
- **External AI Services Layer:** Provides LLM reasoning through Gemini API and Groq LLM API.

## 2.2 Process Flow

The system follows a sequential workflow where each stage depends on the output of the previous stage.



## 2.3 Information Flow

Information flow describes how data moves between different layers and modules.

### Input Information

- User-entered symptoms (text)
- Selected model (Gemini / Groq)
- API key (entered temporarily for session use)

### Processing Information


- Prompts generated from templates
- Context passed between sequential prompts
- Responses received from LLM APIs
- Output parsing and structuring

### Output Information

The final output is shown in the following structured sections:


- Possible Conditions
- Precautions
- General Medication Guidance
- Nutrition Recommendations

### Flow Representation



## 2.4 Components Design

The system consists of multiple components grouped into layers.



## 2.5 Key Design Considerations

The following design considerations were applied during system development:

### 1. Medical Safety

- The system is advisory only.
- The system avoids diagnosis claims.

### 2. Modular Architecture

- Separation into layers improves maintainability.
- Each module can be modified without affecting others.

### 3. Multi-Model Support

- Supports Gemini and Groq models.
- Allows better reliability and response comparison.

### 4. Privacy

- No user medical data is stored permanently.
- API keys are not saved.

## 5. Performance

- System is optimized for single-user prototype usage.
- Response time depends mainly on LLM API latency.

## 2.6 API Catalogue

This section lists the APIs used in the Healthcare Planning Assistant.

### 2.6.1 Google Gemini API

**Purpose:** Generates AI-based healthcare advisory responses.

**Provider:** Google

**Usage in System:** Used as one of the selectable LLM models.

#### Inputs:

- Prompt text
- API key

#### Outputs:

- Generated advisory response in natural language

### 2.6.2 Groq LLM API (Llama Models)

**Purpose:** Generates AI-based healthcare advisory responses using Groq-hosted Llama.

**Type:** External REST API (HTTPS)

**Provider:** Groq

**Usage in System:** Used as an alternative selectable LLM model.

#### Inputs:

- Prompt text
- API key

#### Outputs:

- Generated advisory response in natural language

## 3. Data Design

### 3.1 Data Model

The Healthcare Planning Assistant is currently designed as a prototype system and does not use a persistent database. The system mainly works with temporary runtime data collected from the user and generated by the LLM during execution.

#### Key Data Entities

##### 1. User Input Data

- Symptoms (string / text)
- Selected Model (Gemini / Groq)
- API Key (temporary, session-based)

##### 2. Generated Advisory Data

- Possible Conditions (list of strings)
- Precautions (list of strings)
- Medication Guidance (list of strings)
- Nutrition Recommendations (list of strings)

##### 3. Session Data

- Current request state
- Intermediate LLM responses
- Final structured output

### 3.2 Data Access Mechanism

Since the system does not use a database, all data access is handled through in-memory mechanisms.

#### Data Access Methods

- **User Input Access:** Streamlit form inputs
- **Session Storage:** `st.session_state` for storing intermediate outputs
- **LLM Data Access:** API calls via LangChain integration

- **Output Access:** Structured results displayed through Streamlit components (`st.markdown`, `st.write`, etc.)

## External Data Access

The system communicates with external AI services through HTTPS-based APIs:

- Google Gemini API
- Groq LLM API

## 3.3 Data Retention Policies

The Healthcare Planning Assistant is designed with privacy in mind. Since healthcare-related data is sensitive, the system avoids storing any personal medical data permanently.

### Retention Policy

- User symptom data is stored only during the active session.
- API keys are not stored permanently.
- All generated responses are kept temporarily for display purposes.
- Once the user refreshes or closes the application, session data is cleared automatically.

### Privacy Compliance Consideration

Since no persistent storage is used, the system reduces risks related to:

- patient data leakage
- long-term storage misuse
- compliance requirements

## 3.4 Data Migration

Data migration is not applicable in the current prototype because:

- No database is implemented
- No persistent storage is used
- No historical patient records are stored

## 4. Interfaces

### 4.1 User Interface (UI)

The Healthcare Planning Assistant provides a simple web-based user interface developed using Streamlit. The UI is designed to be minimal, interactive, and easy to use for non-technical users.

#### UI Features

- Symptom input textbox (multi-line)
- Model selection dropdown (Gemini / Groq)
- API key input field (secure input)
- Submit button to start processing
- Output displayed in structured sections:
  - Possible Conditions
  - Precautions
  - General Medication Guidance
  - Nutrition Recommendations
- Error messages for invalid input or missing API keys
- Medical disclaimer displayed in results

#### UI Design Considerations

- Clear section headings for readability
- Minimal user steps (input → submit → output)
- Responsive layout supported by Streamlit

### 4.2 External Interface (LLM APIs)

The system interacts with external AI services using HTTPS-based API calls. These services provide natural language reasoning and advisory responses.

#### External APIs Used

- Google Gemini API
- Groq LLM API (Llama Models)

## 4.3 Hardware Interface

No special hardware interface is required. The system runs on a standard machine capable of executing Python applications.

### Minimum Requirements

- Laptop/PC
- Stable internet connection (for API calls)
- Web browser (Chrome, Edge, Firefox)

## 4.4 Software Interface

The software interfaces used by the system include:

### Frontend

- Streamlit

### Backend / Logic

- Python

### LLM Orchestration

- LangChain

### External Services

- Google Gemini API
- Groq LLM API

## 5. State and Session Management

The Healthcare Planning Assistant is a Streamlit-based web application. Since Streamlit applications rerun the script on each user interaction, the system requires a structured session management approach to preserve user inputs, intermediate LLM outputs, and final results during the active session.

### Session State Elements

The following data is maintained temporarily in the session:

#### 1. User Inputs

- `symptoms`: User-entered symptom description
- `selected_model`: Gemini / Groq
- `api_key`: User API key (not stored permanently)

#### 2. Intermediate Workflow Data

- `conditions_output`
- `precautions_output`
- `medications_output`
- `nutrition_output`

#### 3. Final Output

- `final_structured_result`: Aggregated and formatted response displayed in UI

#### 4. Control Flags

- `is_submitted`: Tracks whether the user has submitted symptoms
- `has_error`: Tracks validation/API failures
- `error_message`: Stores error details for UI display

## Session Lifecycle

### Session Start

A session begins when the user opens the Streamlit application. Default values for session variables are initialized if not already present.

## Session Execution

During processing:

- user inputs are stored in session state
- sequential prompt outputs are stored step-by-step
- the final structured output is stored for rendering

## Session End

The session ends when:

- the user closes the browser tab, or
- the application is refreshed/restarted

## Session Security Considerations

- API keys are only stored in memory during the active session.
- No session data is saved to disk or database.
- The system avoids storing any personally identifiable information (PII).
- All LLM communication happens via secure HTTPS.

## Future Enhancement

In future versions, session management can be extended with:

- user login and authentication
- database-based user history
- encrypted storage for session data
- role-based access control for admin dashboards

## 6. Caching

The Healthcare Planning Assistant uses external LLM APIs (Google Gemini and Groq) for generating responses. Since LLM calls are time-consuming and cost-based, caching is an important design feature to improve performance, reduce repeated API calls, and enhance user experience.

## 7. Non-Functional Requirements

### 7.1 Security Aspects

Since the system deals with health-related inputs, security and privacy are critical even in prototype form. The Healthcare Planning Assistant is designed to minimize risk by avoiding storage of sensitive user information.

#### Key Security Requirements

- **API Key Protection:**  
API keys entered by the user are used only during the active session and are not stored permanently.
- **No Persistent Storage of Patient Data:**  
The system does not store symptoms, health conditions
- **Secure Communication:**  
All communication with external AI services (Gemini API and Groq API) occurs through secure HTTPS.
- **Session-Based Isolation:**  
User data is stored temporarily using Streamlit session state, ensuring that one user's session data is not shared with another.
- **Input Validation:**  
The system validates symptom input to prevent invalid or empty requests and reduce prompt injection risks.
- **Medical Disclaimer and Safety Warnings:**  
The system displays disclaimers to ensure the user understands that the output is advisory and not a medical diagnosis.

#### Security Limitations

- No authentication mechanism is implemented.
- No encryption at rest is required since no persistent storage is used.
- The system is not certified for healthcare regulatory compliance.

### 7.2 Performance Aspects

The system performance is primarily dependent on external LLM API response time. The prototype is optimized for single-user local execution.

## Key Performance Requirements

- **LLM Response Time:**

The expected response time depends on API latency and model complexity. Average response time may range from 5 to 20 seconds depending on the selected model.

- **Sequential Execution Efficiency:**

Since the system runs 4 sequential LLM calls (conditions, precautions, medication, nutrition), the total time is the combined latency of all calls.

- **Session Caching:**

Intermediate results are stored in `st.session_state` to prevent repeated calls during UI reruns.

- **Lightweight Processing:**

Most processing is limited to prompt formatting and response parsing, which is computationally lightweight

## Performance Limitations

- Not optimized for multiple concurrent users.
- No background job queue is used.
- API limits and throttling may affect performance.

## Future Enhancements for Performance

- Combine multiple steps into fewer LLM calls.
- Add response caching with TTL.
- Use async API calls for faster pipeline execution.
- Deploy on cloud infrastructure for scalability.

## 8. References

### 1. Streamlit Documentation

Streamlit Official Docs – UI framework for Python web applications.

### 2. LangChain Documentation

LangChain Official Docs – Framework for LLM orchestration, prompt templates, and chaining.

### 3. Google Gemini API Documentation

Google AI Gemini API – Official documentation for Gemini model integration.

### 4. Groq LLM API Documentation

Groq Cloud API – Official documentation for accessing Llama-based models.

### 5. Python Documentation

Python Official Docs – Core language reference and standard libraries.